

1. Overview

Based on the descriptive and exploratory analysis done in notebook 00_data_understanding, this Python Script will work on preprocessing the data, preparing it so that we can then work on the model training in the future.

2. Data Understanding

2.1 Data Description

This file will use the df_train_transform excel sheet created in the previous notebook: 00_data_understanding

2.2 Import Necessary Libraries

```
In [1]: 1 # pip install category_encoders
```

```
In [2]: 1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 import seaborn as sns
6 from sklearn.preprocessing import OneHotEncoder
7 from category_encoders import TargetEncoder
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.model_selection import train_test_split
10
11
12 import pickle
13 import warnings
14 warnings.filterwarnings('ignore')
15
```

3. Code

3.1 Import the database

```
In [3]: 1 df = pd.read_excel('df_train_transform.xlsx')
        2 df.head()
```

```
Out[3]:
```

	amount_tsh	gps_height	population	basin	region	public_meeting	permit	extraction_typ
0	6000.0	1390	109	lake nyasa	iringa	1.0	0.0	
1	0.0	1399	280	lake victoria	mara	NaN	1.0	
2	25.0	686	250	pangani	manyara	1.0	1.0	
3	0.0	263	58	ruvuma southern coast	mtwara	1.0	1.0	sub
4	0.0	0	0	lake victoria	kagera	1.0	1.0	

3.2 Class Imbalance checking

```
In [4]: 1 # Check class distribution in y_train
        2 print("Class distribution of status_group:")
        3 print(df['status_group'].value_counts(normalize=False))
```

```
Class distribution of status_group:
functional          32259
non functional      22824
functional needs repair  4317
Name: status_group, dtype: int64
```

```
In [5]: 1 # Check class distribution in y_train
        2 print("Class distribution of status_group:")
        3 print(df['status_group'].value_counts(normalize=True))
```

```
Class distribution of status_group:
functional          0.543081
non functional      0.384242
functional needs repair  0.072677
Name: status_group, dtype: float64
```

We decide to group together into a same class functional needs repair and functional. In this way, we have a binary classification problem

```
In [6]: 1 # Replace 'functional needs repair' with 'functional'
2 df['status_group'] = df['status_group'].replace('functional needs repair',
3
4 # Verify changes by checking the class distribution again in y_train and y
5 print("Class distribution in y_train after replacement:")
6 print(df['status_group'].value_counts(normalize=True))
```

Class distribution in y_train after replacement:


functional 0.615758

non functional 0.384242

Name: status_group, dtype: float64

To have further insight, let's do a bar graph representation of the distribution of the target variable

```
In [7]: 1 class_distribution = df['status_group'].value_counts(normalize=True)
2
3 # Plotting the bar plot
4 plt.figure(figsize=(10, 6))
5 sns.barplot(x=class_distribution.index, y=class_distribution.values, palette=
6 plt.xlabel('Status Group')
7 plt.ylabel('Proportion')
8 plt.title('Distribution of Functional and Non-Functional Pumps')
9 plt.ylim(0, 1)
10 plt.show()
```



3.3 Define predictor and target variables

```
In [8]: 1 y = df['status_group']
        2 X = df.drop('status_group', axis=1)
```

3.4 Do a train test split

```
In [9]: 1 # Split the data into training and testing sets
        2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
```

3.5 Dealing with null values

```
In [10]: 1 # For train data
         2 (X_train.isna().sum()/len(df))*100
```

```
Out[10]: amount_tsh          0.000000
         gps_height          0.000000
         population          0.000000
         basin              0.000000
         region             0.000000
         public_meeting      4.526936
         permit             4.106061
         extraction_type_class 0.000000
         management_group    0.000000
         payment_type        0.000000
         quality_group       0.000000
         quantity_group      0.000000
         source_type         0.000000
         waterpoint_type     0.000000
         funder_type         0.000000
         installer_type      0.000000
         scheme_management_grouped 0.000000
         dtype: float64
```

Column 'public_meeting'

```
In [11]: 1 X_train["public_meeting"].value_counts(normalize=True)
```

```
Out[11]: 1.0    0.908813
         0.0    0.091187
         Name: public_meeting, dtype: float64
```

```
In [12]: 1 # Given that the null values are only 6%, Lets replace them with the mode
          2
          3 # Calculate the mode of the 'public_meeting' column
          4 public_meeting_mode = X_train['public_meeting'].mode()[0]
          5
          6 # Fill missing values in 'public_meeting' of X_train with the mode from X_train
          7 X_train['public_meeting'].fillna(public_meeting_mode, inplace=True)
          8
          9 # Fill missing values in 'public_meeting' of X_test with the mode from X_train
         10 X_test['public_meeting'].fillna(public_meeting_mode, inplace=True)
         11
         12 # Convert the 'public_meeting' column to type object in both X_train and X_test
         13 X_train['public_meeting'] = X_train['public_meeting'].astype(object)
         14 X_test['public_meeting'] = X_test['public_meeting'].astype(object)
         15
         16 # Verify if all NA values are filled
         17 print(df['public_meeting'].value_counts(normalize=True))
```

```
1.0    0.909838
0.0     0.090162
Name: public_meeting, dtype: float64
```

```
In [13]: 1 public_meeting_mode
```

```
Out[13]: 1.0
```

Column 'permit'

```
In [14]: 1 df["permit"].value_counts(normalize=True)
```

```
Out[14]: 1.0    0.68955
          0.0    0.31045
          Name: permit, dtype: float64
```

```

In [15]: 1 # Given that the null values are only 5%, Lets replace them with the mode
          2
          3 # Calculate the mode of the 'permit' column
          4 permit_mode = X_train['permit'].mode()[0]
          5
          6 # Fill missing values in 'permit' of X_train with the mode of X_train
          7 X_train['permit'].fillna(permit_mode, inplace=True)
          8
          9 # Fill missing values in 'permit' of X_test with the mode of X_train
         10 X_test['permit'].fillna(permit_mode, inplace=True)
         11
         12 # Convert the 'permit' column to type object in both X_train and X_test
         13 X_train['permit'] = X_train['permit'].astype(object)
         14 X_test['permit'] = X_test['permit'].astype(object)
         15
         16 # Verify if all NA values are filled
         17 print(X_train['permit'].value_counts(normalize=True))

```

1.0 0.704272
0.0 0.295728
Name: permit, dtype: float64

```

In [16]: 1 permit_mode

```

```

Out[16]: 1.0

```

3.6 Doing target encoder on the categorical columns

Let's perform a one hot encoder on the categorical columns that have less than 6 categories

```

In [17]: 1 # Identifying categorical columns
          2 categorical_columns = X_train.select_dtypes(include=['object', 'category'])
          3
          4 # Printing the list of categorical columns
          5 print("Categorical columns in X_train:")
          6 print(categorical_columns)

```

Categorical columns in X_train:
Index(['basin', 'region', 'public_meeting', 'permit', 'extraction_type_classes',
 'management_group', 'payment_type', 'quality_group', 'quantity_group',
 'source_type', 'waterpoint_type', 'funder_type', 'installer_type',
 'scheme_management_grouped'],
 dtype=object)

X_train

Let's do a code to apply one hot encoder on the columns that have less than 6 variables and a target encoder on the columns that have more than 6 variables. The reason why we decide to not apply target encoding to all the columns directly is to avoid overfitting

```
In [18]: 1 # Check if 'y_train' and 'y_test' need to be converted to a numeric type
2 if y_train.dtype == 'object':
3     y_train = y_train.astype('category').cat.codes
4 if y_test.dtype == 'object':
5     y_test = y_test.astype('category').cat.codes
6
7 # Capture categorical columns from X_train for encoding
8 categorical_columns = X_train.select_dtypes(include=['object', 'category'])
9
10 # Initialize encoders
11 target_encoder = TargetEncoder()
12
13 # Encoding the categorical columns in X_train and X_test
14 for col in categorical_columns:
15     if X_train[col].nunique() <= 6:
16         # Apply OneHotEncoder for columns with 6 or fewer unique values
17         X_train = pd.get_dummies(X_train, columns=[col], drop_first=True)
18         X_test = pd.get_dummies(X_test, columns=[col], drop_first=True)
19     else:
20         # Apply TargetEncoder for columns with more than 6 unique values
21         X_train[col] = target_encoder.fit_transform(X_train[col], y_train)
22         X_test[col] = target_encoder.transform(X_test[col])
23         pickle.dump(target_encoder, open(f"model_objects/{col}_target_enc.pkl", 'wb'))
24
25 # Display the DataFrame to check the results
26 X_train.head()
27
```

Out[18]:

	amount_tsh	gps_height	population	basin	region	extraction_type_class	payment_t
3607	50.0	2092	160	0.346722	0.315956	0.300187	0.277
50870	0.0	0	0	0.346722	0.443875	0.309484	0.475
20413	0.0	0	0	0.485901	0.398196	0.805243	0.475
52806	0.0	0	0	0.311216	0.398196	0.300187	0.226
50091	300.0	1023	120	0.432348	0.398697	0.805243	0.306

5 rows × 34 columns

3.7 Dealing with numerical columns

X_train

```
In [19]: 1 # Capture numerical columns
2 numerical_columns = X_train.select_dtypes(include=['int64', 'float64']).columns
3
4 # Initialize the StandardScaler
5 scaler = StandardScaler()
6
7 # Fit and transform the numerical columns
8 scaler.fit(X_train[numerical_columns])
9
10 X_train[numerical_columns] = scaler.transform(X_train[numerical_columns])
11
12 # Save the fitted variables
13 pickle.dump(scaler, open(f"model_objects/numerical_columns_scaler.pkl", 'w'))
14
15 # Display the DataFrame to check the results
16 X_train.head()
```

Out[19]:

	amount_tsh	gps_height	population	basin	region	extraction_type_class	payment
3607	-0.084999	2.053863	-0.041306	-0.540016	-0.633090	-0.521411	-0.89
50870	-0.100621	-0.965049	-0.379739	-0.540016	0.555492	-0.463637	0.7i
20413	-0.100621	-0.965049	-0.379739	1.471270	0.131062	2.617222	0.7i
52806	-0.100621	-0.965049	-0.379739	-1.053126	0.131062	-0.521411	-1.3i
50091	-0.006889	0.511216	-0.125914	0.697368	0.135714	2.617222	-0.6i

5 rows × 34 columns

```
In [20]: 1 X_train[numerical_columns].columns
```

Out[20]: Index(['amount_tsh', 'gps_height', 'population', 'basin', 'region',
'extraction_type_class', 'payment_type', 'source_type',
'waterpoint_type', 'installer_type'],
dtype='object')

```
In [21]: 1 numerical_columns
2
3 print(len(numerical_columns))
```

10

X_test


```
In [22]: 1 X_test[numerical_columns] = scaler.transform(X_test[numerical_columns])
          2
          3 # Display the DataFrame to check the results
          4 X_test.head()
```

Out[22]:

	amount_tsh	gps_height	population	basin	region	extraction_type_class	payment
2980	-0.100621	-0.965049	-0.379739	0.205860	-0.699807	2.617222	1.09
5246	-0.100621	-0.965049	-0.379739	0.205860	1.453840	-0.463637	0.7i
22659	-0.097497	1.452101	-0.066689	-0.540016	-0.633090	-0.521411	-0.89
39888	-0.100621	-0.965049	-0.379739	1.471270	0.131062	-0.463637	0.7i
13361	-0.084999	0.635320	0.117334	-0.540016	0.663779	1.165688	-0.89

5 rows × 34 columns

3.8 Concatenate train on one side and test on the other

```
In [23]: 1 # Concatenate all train
          2 df_train = pd.concat([X_train, y_train], axis=1)
          3
          4 # Concatenate all test
          5 df_test = pd.concat([X_test, y_test], axis=1)
          6
          7 # Create a label column
          8 df_train['is_test'] = 0
          9 df_test['is_test'] = 1
```

3.9 Concatenate everything in one dataframe

```
In [24]: 1 data_processed = pd.concat([df_train,df_test], axis=0)
          2
          3 # Reset index
          4 data_processed = data_processed.reset_index(drop=True)
          5
          6 # Rename column 0 to status_group
          7 data_processed = data_processed.rename(columns={0: 'status_group'})
          8
          9 data_processed
```

Out[24]:

	amount_tsh	gps_height	population	basin	region	extraction_type_class	payment
0	-0.084999	2.053863	-0.041306	-0.540016	-0.633090	-0.521411	-0.89
1	-0.100621	-0.965049	-0.379739	-0.540016	0.555492	-0.463637	0.77
2	-0.100621	-0.965049	-0.379739	1.471270	0.131062	2.617222	0.77
3	-0.100621	-0.965049	-0.379739	-1.053126	0.131062	-0.521411	-1.33
4	-0.006889	0.511216	-0.125914	0.697368	0.135714	2.617222	-0.64
...
59395	-0.038133	1.596408	0.741319	-1.230325	-1.769052	-0.521411	-1.33
59396	0.055600	1.704639	-0.062458	-0.569630	-1.180350	-0.521411	-0.64
59397	-0.100621	-0.965049	-0.379739	0.335579	0.103144	-0.521411	0.77
59398	-0.100621	-0.038596	-0.377623	0.697368	0.135714	-0.521411	0.77
59399	-0.100621	1.098547	-0.377623	-0.569630	0.234762	-0.521411	0.77

59400 rows × 36 columns

4. Export the data

```
In [25]: 1 data_processed.to_excel('df_data_processed.xlsx', index=False)
```